

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv("16_Sleep_health_and_lifestyle_dataset.csv")
df
```

Out[2]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	
...
369	370	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	
370	371	Female	59	Nurse	8.0	9	75	3	Overweight	140/95	
371	372	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	
372	373	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	
373	374	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	

374 rows × 13 columns



```
In [3]: df.head()
```

Out[3]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	77
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85

DATA CLEANING AND DATA PREPROCESSING

In [4]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374 entries, 0 to 373
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   Person ID                            374 non-null    int64
1   Gender                              374 non-null    object
2   Age                                 374 non-null    int64
3   Occupation                          374 non-null    object
4   Sleep Duration                      374 non-null    float64
5   Quality of Sleep                    374 non-null    int64
6   Physical Activity Level              374 non-null    int64
7   Stress Level                        374 non-null    int64
8   BMI Category                        374 non-null    object
9   Blood Pressure                      374 non-null    object
10  Heart Rate                          374 non-null    int64
11  Daily Steps                         374 non-null    int64
12  Sleep Disorder                      374 non-null    object
dtypes: float64(1), int64(7), object(5)
memory usage: 38.1+ KB
```

In [5]:

df.describe()

Out[5]:

	Person ID	Age	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	Heart Rate	Daily Steps
count	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
mean	187.500000	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775	6816.844920
std	108.108742	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676	1617.915679
min	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000	3000.000000
25%	94.250000	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000	5600.000000
50%	187.500000	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000	7000.000000
75%	280.750000	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000	8000.000000
max	374.000000	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000	10000.000000

In [6]:

df.columns

Out[6]: Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration', 'Quality of Sleep', 'Physical Activity Level', 'Stress Level', 'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps', 'Sleep Disorder'], dtype='object')

In [7]:

df1=df.dropna(axis=1)
df1

Out[7]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate	Daily Steps
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	72	10000
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	78	12000
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	78	12000
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	8000
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	8000
...
369	370	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	65	6000
370	371	Female	59	Nurse	8.0	9	75	3	Overweight	140/95	65	6000
371	372	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	65	6000
372	373	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	65	6000
373	374	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	65	6000

374 rows × 13 columns



In [8]:

df1.columns

Out[8]: Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration', 'Quality of Sleep', 'Physical Activity Level', 'Stress Level', 'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps', 'Sleep Disorder'], dtype='object')

In [9]:

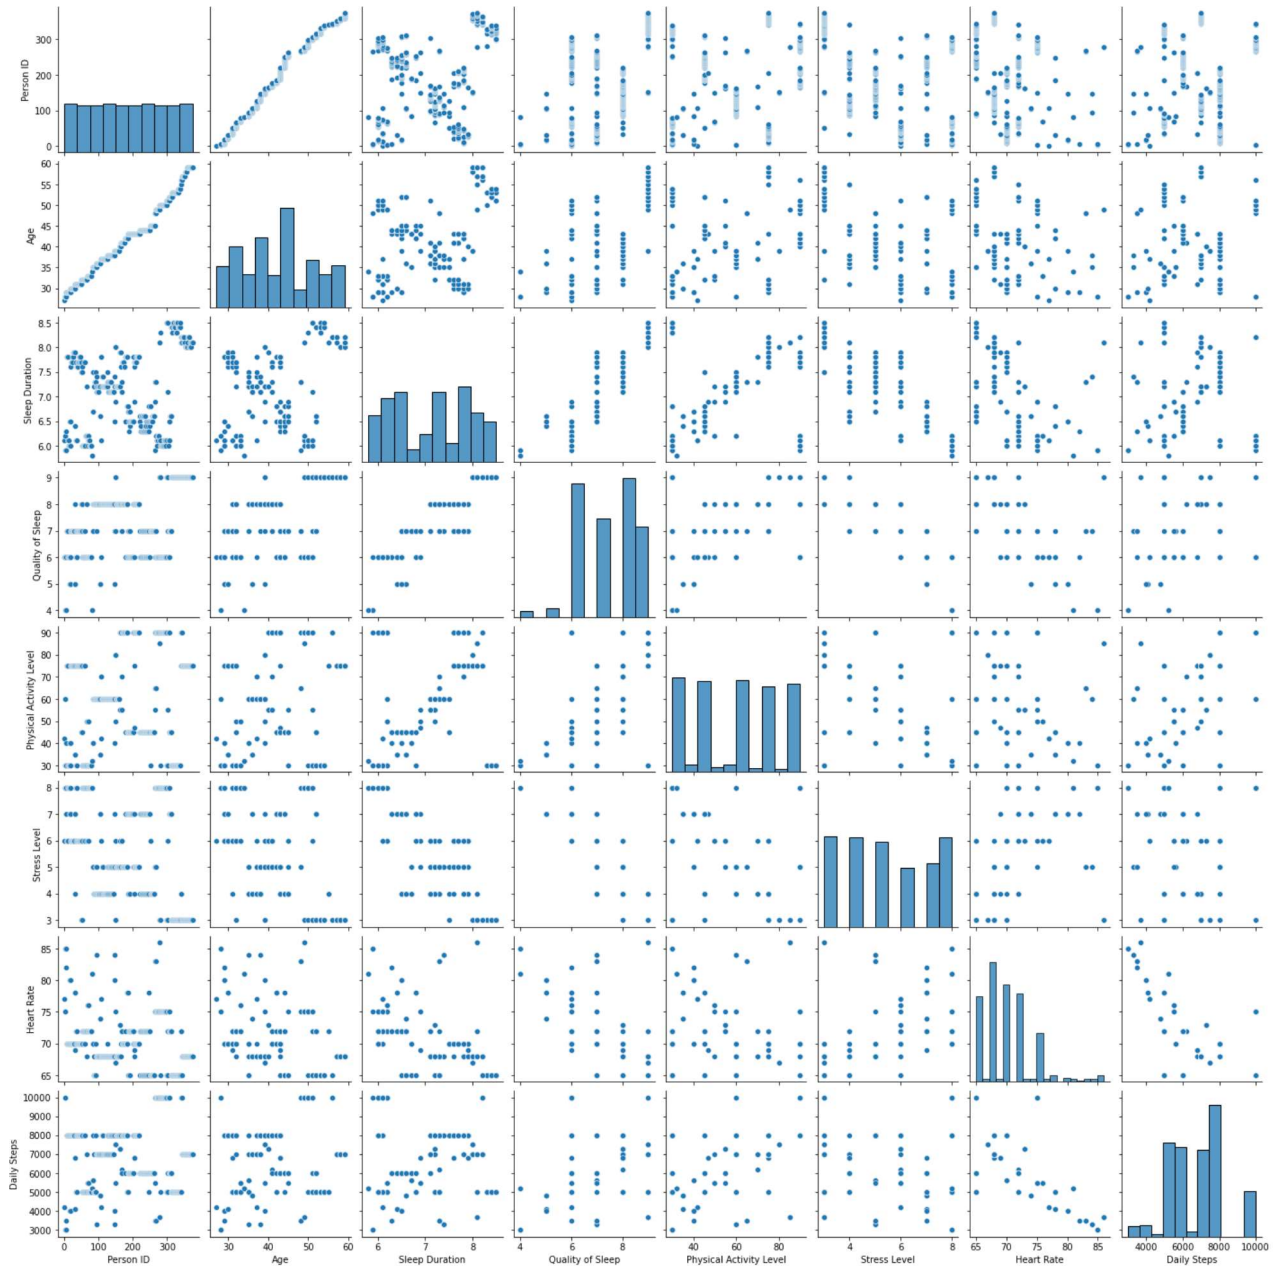
df1=df1[['Person ID', 'Age', 'Sleep Duration', 'Quality of Sleep', 'Physical Activity Level', 'Stress Level',

```
'Heart Rate', 'Daily Steps']]
```

EDA AND VISUALIZATION

```
In [10]: sns.pairplot(df1)
```


```
Out[10]: <seaborn.axisgrid.PairGrid at 0x200e12e7a90>
```



```
In [11]: sns.distplot(df1['Daily Steps'])
```


C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[11]: <AxesSubplot:xlabel='Daily Steps', ylabel='Density'>
```



```
In [12]: sns.heatmap(df1.corr())
```

```
Out[12]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BULDING

```
In [13]: x=df[['Person ID', 'Age', 'Sleep Duration',
               'Quality of Sleep', 'Physical Activity Level', 'Stress Level',
               'Heart Rate']]
          y=df['Daily Steps']
```

```
In [14]: from sklearn.model_selection import train_test_split
          x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [15]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[15]: LinearRegression()

```
In [16]: lr.intercept_
```

Out[16]: 13515.558212288111


```
In [17]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[17]:

	Co-efficient
Person ID	-5.795612
Age	83.052437
Sleep Duration	-314.798189
Quality of Sleep	103.045557
Physical Activity Level	68.085220
Stress Level	480.714306
Heart Rate	-203.517786

```
In [18]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[18]: <matplotlib.collections.PathCollection at 0x200e68fcb20>



ACCURACY

```
In [19]: lr.score(x_test,y_test)
```

Out[19]: 0.7793825557571524

In [20]: `lr.score(x_train,y_train)`

Out[20]: 0.8094049031924386

In [21]: `from sklearn.linear_model import Ridge,Lasso`

In [22]: `rr=Ridge(alpha=10)
rr.fit(x_train,y_train)`

Out[22]: Ridge(alpha=10)

In [23]: `rr.score(x_test,y_test)`

Out[23]: 0.7717218513559383

In [24]: `rr.score(x_train,y_train)`

Out[24]: 0.808892930071658

In [25]: `la=Lasso(alpha=10)
la.fit(x_train,y_train)`

Out[25]: Lasso(alpha=10)

In [26]: `la.score(x_test,y_test)`

Out[26]: 0.7702455753146509

In [27]: `la.score(x_train,y_train)`

Out[27]: 0.8086514574182113

In [28]: `from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)`

Out[28]: ElasticNet()

In [29]: `en.coef_`

Out[29]: array([-1.8956582 , 30.93382906, -125.69893783, -104.2890219 ,
 67.84712099, 306.30222362, -165.08878982])

```
In [30]: en.intercept_
```

```
Out[30]: 13412.620437595991
```

```
In [31]: prediction=en.predict(x_test)
```

```
In [32]: en.score(x_test,y_test)
```

```
Out[32]: 0.741878397881935
```

```
In [33]: from sklearn import metrics  
print(metrics.mean_absolute_error(y_test,prediction))  
print(metrics.mean_squared_error(y_test,prediction))  
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
606.2081947260924  
690759.2012349168  
831.119246098246
```