

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv("11_winequality-red (1).csv")
df
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	qu
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	

1599 rows × 12 columns



```
In [3]: df.head()
```

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	qualit
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	



DATA CLEANING AND DATA PREPROCESSING

In [4]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

In [5]:

df.describe()

Out[5]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.997800	3.511567	0.561349	9.400681	8.651082
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.006798	0.109605	0.007091	0.508113	0.311135
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.996000	3.010000	0.480000	8.000000	5.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.997000	3.310000	0.520000	9.000000	6.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.997500	3.430000	0.540000	9.200000	7.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.998000	3.590000	0.560000	9.400000	8.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	1.000000	4.010000	0.680000	16.000000	10.000000

In [6]:

df.columns

Out[6]:

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
      'pH', 'sulphates', 'alcohol', 'quality'],
      dtype='object')
```

In [7]:

df1=df.dropna(axis=1)
df1

Out[7]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	8.0

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
1	7.8	0.880	0.00	2.6	0.098	25.0	87.0	0.99680	3.20	0.68	9.8	
2	7.8	0.760	0.04	2.3	0.092			0.99700	3.26	0.65	9.8	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
...	
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	

1599 rows × 12 columns



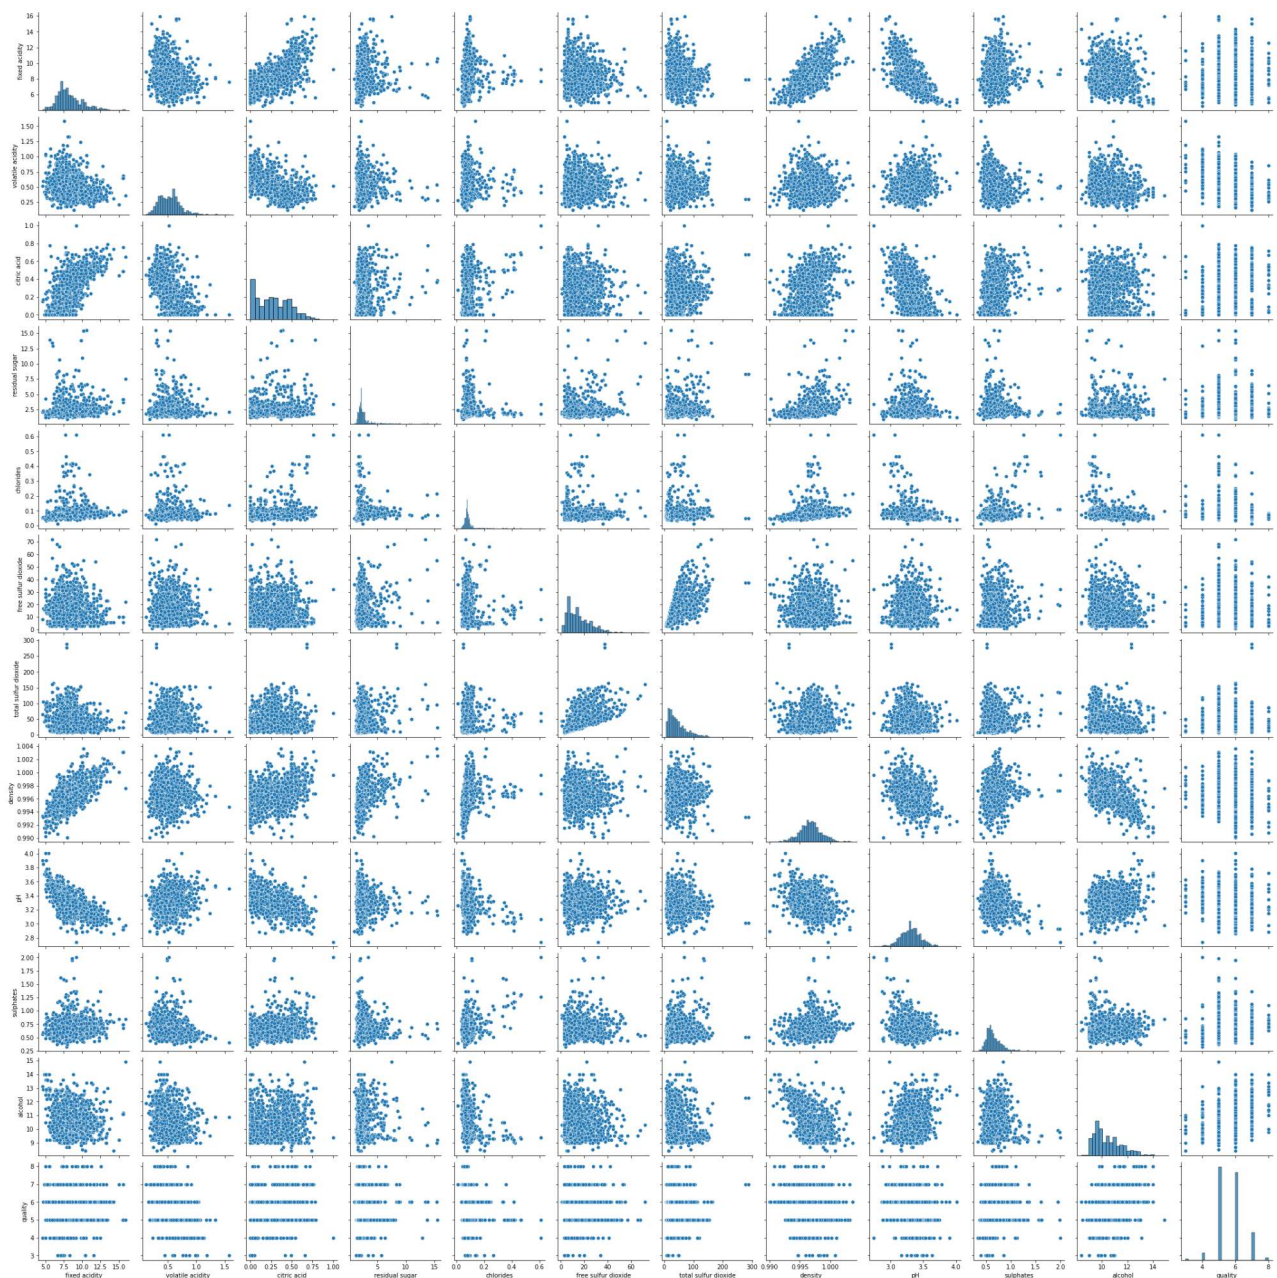
```
In [8]: df1.columns
```

```
Out[8]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
              'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
              'pH', 'sulphates', 'alcohol', 'quality'],  
             dtype='object')
```

EDA AND VISUALIZATION

```
In [9]: sns.pairplot(df1)
```


```
Out[9]: <seaborn.axisgrid.PairGrid at 0x1d2476452e0>
```



```
In [10]: sns.distplot(df1['quality'])
```


C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[10]: <AxesSubplot:xlabel='quality', ylabel='Density'>
```



```
In [11]: sns.heatmap(df1.corr())
```

```
Out[11]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BULDING

```
In [12]: x=df[['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
               'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
               'pH', 'sulphates', 'alcohol']]
          y=df['quality']
```

```
In [13]: from sklearn.model_selection import train_test_split
          x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[14]: LinearRegression()

```
In [15]: lr.intercept_
```

Out[15]: 9.215386337739648

```
In [16]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```


Out[16]:

	Co-efficient
fixed acidity	0.021307
volatile acidity	-0.913780
citric acid	-0.183303
residual sugar	0.029072
chlorides	-1.604660
free sulfur dioxide	0.003804
total sulfur dioxide	-0.003887
density	-5.203263
pH	-0.516513
sulphates	0.864775
alcohol	0.315727

	Co-efficient
fixed acidity	0.021307
volatile acidity	-0.913780
citric acid	-0.183303
residual sugar	0.029072
chlorides	-1.604660
free sulfur dioxide	0.003804
total sulfur dioxide	-0.003887
density	-5.203263
pH	-0.516513
sulphates	0.864775
alcohol	0.315727

```
In [17]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[17]: <matplotlib.collections.PathCollection at 0x1d250f129d0>



ACCURACY

```
In [18]: lr.score(x_test,y_test)
```

```
Out[18]: 0.3368112792092942
```

```
In [19]: lr.score(x_train,y_train)
```

```
Out[19]: 0.3645227919164914
```

```
In [20]: from sklearn.linear_model import Ridge,Lasso
```

```
In [21]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[21]: Ridge(alpha=10)
```

```
In [22]: rr.score(x_test,y_test)
```

```
Out[22]: 0.30858487872183693
```

```
In [23]: rr.score(x_train,y_train)
```

```
Out[23]: 0.354340624245716
```

```
In [24]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[24]: Lasso(alpha=10)
```

```
In [25]: la.score(x_train,y_train)
```

```
Out[25]: 0.0
```

```
In [26]: la.score(x_test,y_test)
```

```
Out[26]: -0.0014521545415211445
```

```
In [27]: from sklearn.linear_model import ElasticNet
          en=ElasticNet()
          en.fit(x_train,y_train)
```

```
Out[27]: ElasticNet()
```

```
In [28]: en.coef_
```

```
Out[28]: array([[ 0.          , -0.          ,  0.          ,  0.          , -0.          ,
                  0.          , -0.00481834, -0.          , -0.          ,  0.          ,
                  0.          ]])
```

```
In [29]: en.intercept_
```

```
Out[29]: 5.850300250095713
```

```
In [30]: prediction=en.predict(x_test)
```

```
In [31]: en.score(x_test,y_test)
```

```
Out[31]: 0.006770986925779265
```

```
In [32]: from sklearn import metrics
          print(metrics.mean_absolute_error(y_test,prediction))
          print(metrics.mean_squared_error(y_test,prediction))
          print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
0.6730626394239672
0.6951913349149353
0.8337813471857807
```